

Metaheuristics for Optimization TP5

Ning, Tientso 24 November, 2019

Introduction

The problem is to create a simple, three-layered, neural-net in order to categorize hand-written digits. The back-propagation (optimization) aspect of the neural-net is done using the Particle Swarm Optimization algorithm (PSO) to minimize loss. The PSO algorithm is a collective-work algorithm that leverages many smaller components in order to yield a collective solution.

Methodology

A neural net is made up of many layers, consisting of nodes and activation functions. For our assignment, the neural-net is a three-layered net, taking an input of an 20×20 image (in flattened form) and outputting a single digit prediction 0 or 1 (corresponding to “3” or “2”).

The first layer of the neural-net takes an input image of 20×20 , and flattens it to a 1×400 vector. It then adds a singular bias value, and passes this 401 vector onto the second layer.

The second layer of the neural-net takes the 401 vector and performs a matrix multiplication with matrix θ_1 (which is part of the $\theta_1 \theta_2$ pair, acting as the weights of the neural-net, and the particle in our PSO). θ_1 has the dimensions 25×401 , turning the resulting matrix into a 1×25 dimension matrix. The sigmoid function is applied in the second layer, along with another bias. The output from layer two is a 1×26 matrix.

The third layer of the neural net applies another matrix multiplication with θ_2 (part of the $\theta_1 \theta_2$ pair). θ_2 is a 26×1 matrix, turning the final result into a 1×1 matrix (or singular value prediction). The sigmoid function is applied, and the output is a prediction value between 0 and 1. During evaluation, the prediction is rounded to the nearest integer value (0 or 1) to count as the prediction, but the value itself before rounding could be considered the “confidence value” of the neural-net.

The PSO algorithm is a process of spawning many particles, each with individual positions and velocities. The particles move in the search-space according to their position and velocity, as well as in relation to the particle with the best fitness. Each particle's movement is defined by it's own position and velocity, and velocities are updated according to the equation:

$$\mathbf{v}_{t+1}^i = \omega \mathbf{v}_t^i + c_1 r_1 (\mathbf{b}_t^i - \mathbf{s}_t^i) + c_2 r_2 (\mathbf{b}_t^G - \mathbf{s}_t^i)$$

The intuition is that each particle will move according to trying to move towards the particle with the best fitness amongst the group, as well as move towards where it had the best fitness itself. These two guiding concepts are established variables and defined as the cognitive and social affinity of each particle (defined in our assignment as c1 and c2). There is also an established Inertia Constant, which acts to help prevent abrupt spikes in velocity. These values assist in making sure that the particles are able to properly explore the search space as well as exploit the search space for an optimal solution. The fitness function is provided as:

$$J(\Theta_1, \Theta_2) := \frac{1}{m} \sum_{k=1}^m (y_k - h_{\Theta_1, \Theta_2}(x_k))^2$$

The set up of the problem establishes the two matrices theta1 and theta2 pair, as a single particle in our PSO. Thus, the problem is about optimizing to see which particle (set of weights) gives us the proper prediction results (learning a set of weights) over the set of images in our dataset.

Results

Run	Fitness Value	Prediction Error
1	0.0020812038887232927	0.0
2	0.0018613057836248423	0.0
3	0.0016785848636784756	0.0
4	0.0032176760351827504	0.0
5	0.0020581811150554400	0.0
6	0.0052839376679271900	0.05
7	0.0024315360392561300	0.0
8	0.0019122750124164996	0.0

Run	Fitness Value	Prediction Error
9	0.0005050333698191652	0.0
10	0.004716182299268048	0.0

(Table 1: Table of Fitness Values and Prediction Errors)



Figure 1: Fitness over Iterations

Discussion

The results show that our algorithm is able to achieve an average fitness value of 0.0025745916, with an average prediction error of 0.05, meaning that the PSO algorithm is able to optimize the fitness of the function by minimizing prediction error, and find the weight matrices theta1 and theta2 to allow the neural-net to achieve its prediction accuracy.

It is important to note, that according to the assignment instructions, we are only making our predictions on the testing set, which explains why our prediction error is 0.0 for all runs. It can be expected that on images we have not trained on, that our error would be larger. However, that is beyond the scope of the assignment and is not a consideration for this time. Additionally, our fitness values indicate that our PSO algorithm is indeed optimizing the weight matrices, as the fitness is close to 0.0.

We chose 20 particles at 200 iterations since those values gave the best results on our search space, and gave us a reasonable run-time. As the number of particles increased, we see an improvement in the exploration of the algorithm allowing us to better converge to an optimal solution. However, as the number of particle increases, we can also expect an increase in run-time. The number of particles help us explore the search space, meaning that an inadequate number of particles could lead to seeing unstable results, having less of a guarantee that the algorithm will converge on the global-optimum. As seen in the graph, as the number of iterations increased, the algorithm was able to better exploit the search space and converge to the global-optimum, allowing our algorithm to perform well. At 200 iterations, we can see that the fitness is well-close to zero, which is close to optimality.

Additionally, the velocity cutoff, V_{max} , is critical to the ability to provide good results since a velocity that is too high would cause convergence issues as particles' positions are determined by adding the velocity to their current position, meaning a velocity that is too high would cause particles to alter its position too much too quickly for converging on a good solution. It is also important to note that we need to restrain the particle to be contained within the search space, which was set to be between -1 and 1 for our assignment. A velocity cutoff of 0.1 was chosen for our assignment, which is one-tenth of the domain of the search-space.

Conclusion

In conclusion, our PSO algorithm was able to optimize the θ_1 θ_2 matrix pair, allowing our neural-net to make correct predictions on the image dataset.